

AI-Verse 3D Explorer:

A Tactile, Interactive 3D Platform for Visualizing and Understanding Neural Network Architectures

Project Documentation Report

WebGL & Three.js Neural Architecture Visualization System

V.V.S.M.Pavanateja

Repository: github.com/VemuriPavanateja2007/Study-Related-Projects/ai-verse-3d-explorer

July 2026

Abstract

The application of interactive three-dimensional (3D) visualization to explain internal computational processes of artificial intelligence (AI) systems is an emerging pedagogical approach for demystifying otherwise abstract mathematical operations. This report documents the design, architecture, and functionality of AI-Verse 3D Explorer, a browser-based, WebGL/Three.js-powered educational platform that renders four canonical neural network architectures — the Deep Neural Network (DNN) with backpropagation, the Convolutional Neural Network (CNN), the Transformer self-attention mechanism, and the Large Language Model (LLM) autoregressive generation engine — as manipulable, rotatable, tactile 3D scenes. Users may orbit, pan, and zoom each architecture, inspect per-layer mathematics and parameters through a technical inspector panel, and export a four-page PDF architecture report summarizing the explored concept. This work discusses the platform's terminology and conceptual framework, its module-by-module functionality, its interaction and rendering design, its educational value, and its current limitations, drawing structural inspiration from established conventions in academic technical reporting.

Keywords: neural network visualization, WebGL, Three.js, interactive learning, deep neural networks, convolutional neural networks, self-attention, large language models, explainable AI, educational technology.

1. Introduction

Understanding how modern artificial intelligence systems operate internally is a persistent challenge for students, practitioners, and even seasoned engineers. Deep learning architectures involve high-dimensional linear algebra, iterative gradient-based optimization, and layered representations that are difficult to communicate through static diagrams or textual description alone. Traditional teaching materials — whiteboard diagrams, two-dimensional plots of neurons and edges, or slide-based explanations of matrix operations — frequently fail to convey the depth, scale, and dynamic behavior of these systems.

AI-Verse 3D Explorer addresses this gap by translating four of the most consequential architectures in modern AI — the fully connected Deep Neural Network (DNN), the Convolutional Neural Network (CNN), the Transformer's self-attention mechanism, and the Large Language Model (LLM) generation engine — into interactive, spatial, three-dimensional scenes that a learner can rotate, pan, zoom into, and inspect. Rather than reading about forward propagation and backpropagation, a learner can watch translucent 3D nodes pulse with a reverse gradient signal. Rather than imagining a convolutional filter sliding over an image tensor, a learner can explode a stack of feature-map layers apart and watch a 3D kernel scan across them.

This report documents the project's motivation, conceptual framework, module-level functionality, implementation, and educational significance. Its structure intentionally follows the conventions of an academic technical paper — introduction, related work and terminology, conceptual framework, detailed module descriptions with figures, discussion, limitations, and conclusion — to provide a rigorous and complete reference document for the project.

1.1 Motivation and Problem Statement

Machine learning curricula typically separate mathematical theory (matrix calculus, gradient descent, softmax normalization) from visual intuition, forcing learners to mentally reconstruct spatial relationships that are never directly shown. This separation is most acute for architectures whose behavior is inherently multi-dimensional: a CNN's hierarchy of feature maps, a Transformer's pairwise attention scores across every token in a sequence, and an LLM's sliding context window across a token stream. AI-Verse 3D Explorer was conceived to close this gap by providing a single tactile workspace in which each of these architectures is rendered as an explorable 3D object rather than a static equation.

1.2 Scope of this Document

- Section 2 reviews the terminology and conceptual distinctions between the architectures covered by the platform.
- Section 3 presents the system's overall framework: platform structure, navigation flow, and technology stack.

- Section 4 documents each of the four visualization modules individually, including annotated figures captured directly from the running application.
- Section 5 describes the technical inspector panel and the PDF export feature shared across all modules.
- Section 6 discusses the pedagogical value and design rationale of the platform.
- Section 7 discusses current limitations and proposes directions for future work.
- Section 8 concludes the report.

2. Background and Terminology

Before describing the platform itself, it is useful to briefly delineate the four architectural families it visualizes, since each demands a distinct spatial metaphor.

2.1 Deep Neural Networks and Backpropagation

A Deep Neural Network (DNN) is a fully connected, feedforward architecture that maps high-dimensional inputs through successive non-linear hidden layers to produce an output. Forward propagation computes a prediction by passing signals from the input layer through each hidden layer to the output layer; backpropagation then computes the reverse error signal, using the chain rule to propagate gradients backward through the network and update synaptic weights. This bidirectional signal flow — forward prediction versus backward gradient pulse — is the central concept the DNN module of the platform aims to make visible.

2.2 Convolutional Neural Networks

A Convolutional Neural Network (CNN) is a spatial processing architecture that applies sliding two-dimensional (and, in this platform's visualization, three-dimensional) kernels across an input tensor to extract hierarchical visual features. Early layers typically detect low-level edges, middle layers detect textures and motifs, and later layers assemble high-level object-part compositions. This hierarchy is naturally suited to a layered 3D representation in which each convolutional stage is rendered as a distinct slice.

2.3 Transformers and Self-Attention

The Transformer architecture's foundational mechanism is self-attention, which computes pairwise contextual relationships between every pair of tokens in an input sequence simultaneously, rather than processing tokens strictly in order. Each token's query vector is compared against every other token's key vector to produce an attention score, which determines how much influence each token exerts on the representation of the token being computed. Because these pairwise scores form a dense two-dimensional matrix, they can be rendered as an elevation surface — a 3D 'landscape' where higher attention scores correspond to higher peaks.

2.4 Large Language Models and Autoregressive Generation

A Large Language Model (LLM) built on stacked Transformer blocks is fundamentally an autoregressive token generator: given a sequence of prior tokens within a bounded context window, it evaluates a probability distribution (via a softmax over output logits) over the vocabulary and samples the next token, appending it to the sequence before repeating the process. This project visualizes both the layered transformer embedding structure and the sliding context window with its resulting candidate-token probability distribution.

3. System Framework and Navigation Flow

AI-Verse 3D Explorer is structured as a three-stage browser experience: a landing page that introduces the platform and its four modules, a concept-selection stage where the user picks (or types) an architecture to explore, and an interactive 3D workspace stage where the chosen architecture is rendered and can be inspected in detail. This staged flow keeps the cognitive load low: a learner is never asked to configure a 3D scene manually, only to choose what to explore and then observe and interact.

3.1 Landing Page

The landing page (Figure 1) presents the platform's identity — 'AI-Verse 3D Explorer: Tactile Neural Architecture Platform' — alongside its central value proposition, 'Touch, Rotate, and Explore the Architecture of Artificial Intelligence.' Beneath the headline, a short description sets learner expectations: the ability to inspect gradient backpropagation, explode CNN feature maps, visualize Transformer 'Attention Cities,' and inspect LLM sliding context windows. A prominent call-to-action button, 'Start Exploring Workspace,' transitions the user to the concept-selection stage. Four preview cards along the bottom of the page summarize each of the platform's four modules: DNN & Backprop, CNN Feature Maps, Attention City, and LLM Helix & PDF, each annotated with a one-line description of its signature interaction mechanic.




Figure 1. Landing page of AI-Verse 3D Explorer, showing the platform title, value proposition, and the four module preview cards.

3.2 Concept Selection Stage

The second stage, 'Choose an Architecture to Explore' (Figure 2), presents the same four architectures as selectable cards — Deep Neural Network (DNN), Computer Vision (CNN), Transformers & NLP, and Large Language Models (LLMs) — each tagged with its research category (Neural Core, Computer Vision, Attention Mechanism, Generative AI) and a short list of the signature interactive elements it contains (for example, 'Forward Signal,' 'Reverse Gradient Pulse,' and 'Activation Functions' for the DNN card). In addition to the four preset cards, a free-text search field allows a user to type any custom AI concept (for example, 'Multi-Head Attention' or 'Mixture of Experts') and generate a tailored workspace, extending the platform beyond its four built-in presets.




Figure 2. Concept-selection stage, displaying the four preset architecture cards and the custom-concept search field.

3.3 Common Workspace Layout

Once a concept is selected, the user enters a common workspace layout shared by all four modules (visible in Figures 3 through 6). This layout consists of: a top navigation bar with a breadcrumb back to 'Concepts,' the current architecture's title and category tag, camera-view shortcuts (Perspective, Top, Front, Side), a Pause/Play control for any animated signal, and a 'Download PDF' shortcut; a large central 3D canvas rendered with orbit controls (left-drag to rotate, right-drag to pan, scroll to zoom); and a right-hand 'Technical Inspector' panel with four tabs — Overview, Layers, Math, and Params — plus a persistent 'Download 4-Page PDF Architecture Report' button at the panel's base. This consistent scaffolding means that once a user learns to operate one module, all remaining modules are immediately navigable.

4. Module-by-Module Documentation

This section documents each of the platform's four core visualization modules in turn, describing the specific 3D metaphor used, the interactive elements available, and the conceptual overview text presented to the learner.

4.1 Module 1 — Deep Neural Network (DNN) and Backpropagation

The DNN module (Figure 3) renders a four-layer fully connected network — an input layer, two hidden layers, and an output layer — as columns of translucent colored spheres (nodes) connected by thin line edges representing synaptic weights. The Core Concept Overview panel explains the architecture as 'a fully connected feedforward architecture mapping high-dimensional inputs through hidden non-linear layers,' noting that forward propagation computes predictions while reverse error signals backpropagate gradients using the chain rule.

Three key architectural pillars are surfaced in the inspector: (1) bidirectional signal flow, contrasting the forward prediction pulse against the backward weight-update gradient pulse; (2) interactive synaptic weights, allowing a user to click individual nodes to inspect activation values, bias terms, and local gradients; and (3) an adjustable learning-rate control (denoted η , eta) that lets the learner observe how gradient magnitude changes in real time as the rate is tuned. The top toolbar additionally exposes a dedicated 'Backprop' toggle and camera presets labeled PER (perspective), TOP, FRO (front), and SID (side).




Figure 3. The Deep Neural Network & Backpropagation module, showing the input, hidden, and output layers as a 3D node graph with the Technical Inspector's Overview tab open.

4.2 Module 2 — Convolutional Neural Network (CNN) Feature Decomposition

The CNN module (Figure 4) represents the network as a vertically stacked sandwich of 3D feature-map layers, beginning with the input image tensor at the top and narrowing through successive convolutional and pooling stages down to a global pooling and fully connected (FC) layer at the base. Labeled stages visible in the workspace include 'Conv2D Layer 1 + ReLU (Edge Detection),' 'Conv2D Layer 2 + MaxPool (Texture & Motifs),' and 'Conv2D Layer 3 (High-Level Part Detector),' culminating in the 'Global Pooling & FC Layer.' A small annotated kernel indicator ('Kernel: 3x3 Conv (3x3)') marks the active filter size.

The inspector describes the CNN as 'a spatial processing network utilizing 2D/3D sliding kernels to extract spatial visual hierarchies — translating raw pixel grids into low-level edges, mid-level textures, and high-level object part compositions.' Its three architectural pillars are: (1) an explodable 3D layer sandwich, in which an Explode Distance slider physically separates the stacked feature slices for closer inspection; (2) a 3D scanning filter kernel that animates a box scanning across feature layers in real time with a kernel-weight overlay; and (3) the hierarchical visual breakdown from edge maps up to high-level classification.




Figure 4. The Convolutional Neural Network Feature Decomposition module, showing the exploded 3D stack of convolutional and pooling layers from input tensor to the global pooling/FC layer.

4.3 Module 3 — Transformers and Self-Attention ("Attention City")

The Transformer module (Figure 5), branded within the platform as 'Attention City,' visualizes self-attention as a topological 3D heightmap: a grid of spike-like elevation peaks whose height encodes the attention score between token pairs. Along the foreground, the input sequence — 'The animal didn't cross the street because it was too tired' — is laid out as a row of individually labeled 3D token blocks. Selecting the token 'it' reveals percentage-labeled attention weights (for example, 80%, 10%, and 10%) connecting it via glowing curved (Bezier) lines to the tokens it attends to most strongly, most notably 'animal.'

The inspector overview explains this as 'the foundational self-attention mechanism that computes pairwise contextual relationships between all tokens in a sequence simultaneously, mapping query-key interactions into a 3D topological heightmap.' Its three pillars are: (1) the 3D Attention City itself, in which attention scores form glowing elevation peaks on a 3D grid; (2) token-hover Bezier connections, letting a user hover over any token to trace curved connection lines to the tokens most relevant to it; and (3) a Multi-Head Switcher, letting the learner toggle between attention heads (H1, H2, and beyond) to compare the distinct relational patterns each head has learned.




Figure 5. The Transformer Self-Attention ('Attention City') module, displaying the 3D attention heightmap, the tokenized input sequence, and the percentage-weighted attention connections from the token 'it'.

4.4 Module 4 — Large Language Model (LLM) Generation Engine

The final module (Figure 6) visualizes an LLM's autoregressive generation process through two coupled 3D elements: on the left, a '3D Transformer Embedding Latent Helix' — a double-helix arrangement of teal and violet spheres representing stacked transformer token/positional embeddings connected across turns of the helix; and on the right, a bounded 'Context Window Box (Length = 8)' containing a bar chart of candidate next-token probabilities (visible values include 51%, 29%, 12%, 5%, 3%, and 1% in descending order), representing the softmax output-logit distribution the model samples from to choose the next token.

The inspector describes the module as 'an autoregressive token generator that processes prompt tokens within a sliding 3D Context Window box, evaluating output logit probability distributions to predict and sample the next token.' Its three pillars are: (1) the 3D double-helix matrix as a visual representation of stacked transformer representations; (2) the sliding context-window box, a bounding volume enclosing the currently active sequence tokens; and (3) interactive temperature and top-p controls, which let a learner watch the candidate-token probability bars shift height in real time as sampling hyperparameters are changed.




Figure 6. The Large Language Model Generation Engine module, showing the 3D embedding helix alongside the sliding context window and next-token probability bar chart.

5. The Technical Inspector Panel and PDF Export

Every module in AI-Verse 3D Explorer shares a common right-hand 'Technical Inspector' panel, organized into four tabs: Overview, Layers, Math, and Params. The Overview tab (shown in Figures 3–6) presents the Core Concept Overview paragraph and the three Key Architectural Pillars described above for each module. The Layers, Math, and Params tabs (not expanded in the captured figures) are designed to progressively surface deeper technical detail — respectively, the per-layer structural breakdown of the architecture, the underlying closed-form mathematical expressions (e.g., the chain-rule gradient update for the DNN, the convolution operation for the CNN, the scaled dot-product attention formula for the Transformer, and the softmax/temperature sampling equations for the LLM), and the adjustable numeric hyperparameters (such as learning rate, kernel size, number of attention heads, temperature, and top-p) that drive each module's live simulation.

5.1 PDF Architecture Report Export

Beneath the four inspector tabs, every module exposes a 'Download 4-Page PDF Architecture Report' action, described in the interface as capturing a snapshot of the current 3D canvas and exporting the accompanying mathematical specifications. This feature allows a learner or instructor to generate a static, shareable four-page reference document for any of the four architectures (or a custom concept) directly from the live 3D exploration session, bridging the platform's interactive mode with a conventional, portable study artifact.

5.2 Camera and Simulation Controls

Across all modules, the top toolbar exposes four camera-view shortcuts — PER (perspective), TOP, FRO (front), and SID (side) — allowing a learner to snap to a standard viewing angle without manually re-orienting the 3D scene via mouse drag. A Pause/Play control freezes or resumes any ongoing signal animation (for instance, the DNN's forward/backward pulse or the LLM's token-by-token generation), which is particularly useful when an instructor wishes to hold a frame in place for classroom discussion. A persistent status bar at the bottom-left of the canvas ('3D OrbitControls Active • Left Drag: Rotate | Right Drag: Pan | Scroll: Zoom') reinforces the interaction model for first-time users.

6. Discussion: Educational Value and Design Rationale

AI-Verse 3D Explorer's central design decision — to give every architecture family its own bespoke spatial metaphor rather than a single generic diagram type — is what allows the platform to convey conceptually distinct ideas without conflating them. The DNN module's node-and-edge graph is the natural representation for a fully connected network's weighted connections; the CNN module's explodable layer sandwich is the natural representation for a strictly hierarchical, spatially local feature pipeline; the Transformer module's heightmap and Bezier connections are the natural representation for a fully pairwise, non-local attention mechanism; and the LLM

module's helix-plus-context-window pairing separates the model's static learned representation (the helix) from its dynamic, per-step generation process (the context window and probability bars).

This differentiation matters pedagogically: learners frequently conflate CNNs and DNNs, or assume that a Transformer's attention mechanism behaves like a CNN's local convolution. By assigning each architecture a distinct, purpose-built 3D visual language, the platform makes these structural differences immediately visible rather than something that must be inferred from equations. The consistent workspace scaffolding (breadcrumb, camera shortcuts, Technical Inspector panel, and PDF export) further ensures that this differentiation in content does not come at the cost of a fragmented user experience — once a learner is oriented in one module, that orientation transfers directly to the next.

The inclusion of a free-text custom-concept field on the concept-selection stage (Section 3.2) also signals an intended extensibility beyond the four built-in architectures, suggesting the underlying rendering framework is designed to generate tailored 3D workspaces for arbitrary AI concepts such as Multi-Head Attention or Mixture-of-Experts routing, rather than being hard-coded to exactly four scenes.

7. Limitations and Future Work

As a conceptual and interactive teaching tool rather than a live model-training environment, AI-Verse 3D Explorer's visualizations are illustrative rather than computed from a real, trainable network instance; the node activations, attention scores, and probability distributions shown are representative examples rather than outputs of an underlying live model executing on real data. Empirical validation of the platform's educational effectiveness — for example, through pre/post learner assessments — has not yet been conducted and represents a natural next step, mirroring the kind of iterative, field-validated development recommended for conceptual frameworks in the broader AI/ML literature.

- Extend the four built-in presets with additional architectures (e.g., Recurrent Neural Networks, Mixture-of-Experts, Diffusion Models) using the same bespoke-metaphor design philosophy.
- Wire the visualizations to a real, lightweight backend model so that node activations, attention scores, and token probabilities reflect genuine forward passes on user-supplied input rather than illustrative placeholder values.
- Expand the Layers, Math, and Params inspector tabs with fully worked derivations and adjustable live hyperparameters for every module, not only the DNN's learning rate and the LLM's temperature/top-p.
- Conduct structured user studies with students and instructors to quantify learning-outcome improvements relative to traditional static-diagram instruction.
- Support collaborative or classroom modes, allowing an instructor's camera view and paused state to be mirrored to student devices in real time.

8. Conclusion

AI-Verse 3D Explorer demonstrates that four of the most influential neural network architectures in contemporary AI — the Deep Neural Network, the Convolutional Neural Network, the Transformer's self-attention mechanism, and the Large Language Model's autoregressive generation engine — can each be given a distinct, tactile, three-dimensional visual language that makes their internal computational behavior directly explorable rather than merely described. Through a consistent three-stage flow (landing page, concept selection, and interactive 3D workspace) and a shared Technical Inspector and PDF-export scaffold, the platform balances architecture-specific visual differentiation with a uniform, learnable interaction model. While current visualizations are illustrative rather than computed from a live trainable model, the platform's underlying framework — a bespoke 3D metaphor per architecture family, layered technical detail via the inspector panel, and exportable static reports — establishes a solid and extensible foundation for interactive, spatial AI education, with clear directions for empirical validation and further architectural coverage in future work.

References

- [1] Kühl, N., Goutier, M., Hirt, R., & Satzger, G. (2019). Machine Learning in Artificial Intelligence: Towards a Common Understanding. Proceedings of the 52nd Hawaii International Conference on System Sciences (HICSS). <https://hdl.handle.net/10125/59960>
- [2] Russell, S. J., & Norvig, P. (2015). Artificial Intelligence: A Modern Approach (3rd ed.).
- [3] LeCun, Y., Bengio, Y., & Hinton, G. E. (2015). Deep learning. *Nature*.
- [4] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems (NeurIPS)*.
- [5] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep Residual Learning for Image Recognition. *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- [6] Vemuri, P. A. (2026). AI-Verse 3D Explorer [Software]. GitHub repository: github.com/VemuriPavanateja2007/Study-Related-Projects/tree/main/ai-verse-3d-explorer
- [7] Three.js Documentation. WebGL 3D Library. <https://threejs.org/docs/>